

Operationalizing an AWS ML Project

Dog Breed Image Classification using AWS SageMaker and PyTorch

In this project, several important tools and features of AWS are used to adjust, improve, configure, and prepare a ML model for production deployment. The project contains code that trains and deploys an image classification model on AWS Sagemaker.


The project is performed following the next steps:

- Step 1: Training and deployment on AWS Sagemaker.
- Step 2: Training using EC2.
- Step 3: Lambda function setup for inference.
- Step 4: Lambda security setup and testing.
- Step 5: Lambda concurrency setup and endpoint auto-scaling.

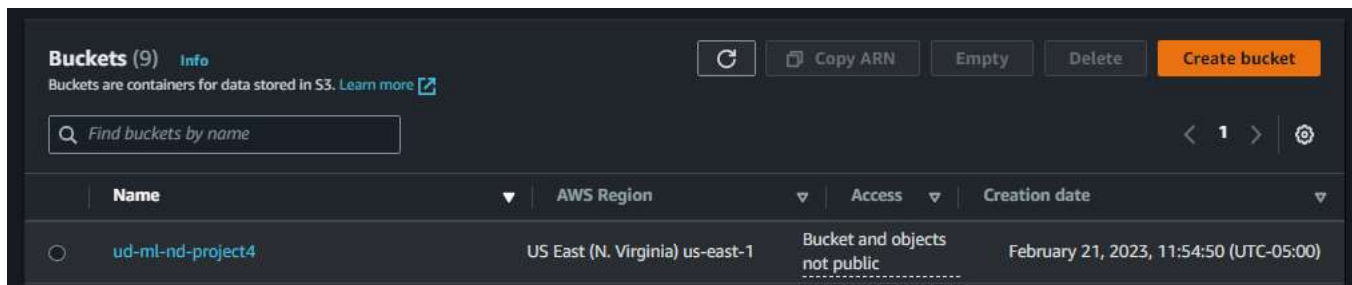
Step 1: Training and deployment on AWS Sagemaker

The training process is performed using a Notebook instance on Sagemaker. A small instance type such as `m1.t3.medium` (2 vCPU and 6 GiB of memory) is selected for the project, since this meets the requirements for running the `train_and_deploy-solution.ipynb` notebook at the lowest cost of \$ 0.05 per hour.

The following image shows the notebook instance created:



A S3 bucket is set up in order to storage the training data:



The model is trained on both single and multiple instances using a `m1.g4dn.xlarge` instance type. For practical purposes, the training halts at the first epoch. Then, the models from single instance and multiple instance training jobs are deployed, as follows:

- Single-instance training:

Log streams (52)

dog-pytorch-2023-02-21-17-57-51-682

1 match

☐ Exact match

☐ Show expired

Info

<

1

>

☐

Log stream

Last event time

☐

dog-pytorch-2023-02-21-17-57-51-682/algo-1-1677002337

2023-02-21 13:00:21 (UTC-05:00)

Amazon SageMaker

Endpoints

Endpoints

Update endpoint

Actions

Create endpoint

Q

Search endpoints

<

1

>

	Name	ARN	Creation time	Status	Last updated
<div></div>	pytorch-inference-2023-02-21-18-21-09-385	arn:aws:sagemaker:us-east-1:076107025568:endpoint/pytorch-inference-2023-02-21-18-21-09-385	Feb 21, 2023 18:21 UTC	<div><div></div>InService</div>	Feb 21, 2023 18:23 UTC

- Multi-instance training:

Log streams

Metric filters

Subscription filters

Contributor Insights

Tags

Data protection - new

Log streams (55)

Delete

Create log stream

Search all log streams

Q dog-pytorch-2023-02-21-18-56-16-972

X

3 matches

☐ Exact match

☐ Show expired

Info

<

1

>

<input type="checkbox"/>	Log stream	Last event time
<input type="checkbox"/>	dog-pytorch-2023-02-21-18-56-16-972/algo-3-1677005845	2023-02-21 13:58:46 (UTC-05:00)
<input type="checkbox"/>	dog-pytorch-2023-02-21-18-56-16-972/algo-2-1677005845	2023-02-21 13:58:46 (UTC-05:00)
<input type="checkbox"/>	dog-pytorch-2023-02-21-18-56-16-972/algo-1-1677005845	2023-02-21 13:58:45 (UTC-05:00)

Amazon SageMaker

>

Endpoints

Endpoints

Update endpoint

Actions ▾

Create endpoint

Q

Search endpoints

<

1

>

	Name ▾	ARN	Creation time ▾	Status ▾	Last updated
<div></div>	pytorch-inference-2023-02-21-19-17-19-046	arn:aws:sagemaker:us-east-1:076107025568:endpoint/pytorch-inference-2023-02-21-19-17-19-046	Feb 21, 2023 19:17 UTC	<div></div> InService	Feb 21, 2023 19:19 UTC

Step 2: Training on EC2

An alternative for training a model is using an EC2 instance instead of Sagemaker. In this scenario, a script adapted from the notebook is necessary. The instance type chosen for training the model on EC2 is the same `m1.g4dn.xlarge`, since it meets the computing requirements at a lower cost for processing and training. As said before, for practical purposes, the training will halt at the first epoch.

The following images shows the EC2 instance and the terminal running the `ec2train1.py` script for training the model.

Instances (1) Info

Connect

Instance state ▾

Actions ▾

Launch instance

Find instance by attribute or tag (case-sensitive)

< 1

<input type="checkbox"/>	Name ▾	Instance ID	Instance state ▾	Instance type ▾	Status check	Alarm status	Availability Zone ▾
<input type="checkbox"/>	ud-ml-nd-p4	i-08e79fde554664ae6	<div><div></div>Running</div>	g4dn.xlarge	<div><div></div>2/2 checks passed</div>	No alarms +	us-east-1a

```
Amazon SageMaker S3 EC2
[root@ip-172-31-93-19 ~]# conda env list
# conda environments:
#
base                  *  /opt/conda
pytorch               /opt/conda/envs/pytorch

[root@ip-172-31-93-19 ~]# source activate pytorch
(pytorch) [root@ip-172-31-93-19 ~]# ls
dogImages.zip  solution.py  TrainedModels
(pytorch) [root@ip-172-31-93-19 ~]# python solution.py
/opt/conda/envs/pytorch/lib/python3.9/site-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and will be removed in 0.15, please use 'weights' instead.
  warnings.warn(
/opt/conda/envs/pytorch/lib/python3.9/site-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or 'None' for 'weights' are deprecated since 0.13 and will be removed in 0.15. The current behavior is equivalent to passing 'weights=ResNet50_Weights.IMAGENET1K_V1'. You can also use 'weights=ResNet50_Weights.DEFAULT' to get the most up-to-date weights.
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet50-0676ba61.pth" to /root/.cache/torch/hub/checkpoints/resnet50-0676ba61.pth | 97.8M/97.8M [00:00<00:00, 329MB/s]
100%
Starting Model Training
saved
(pytorch) [root@ip-172-31-93-19 ~]# ls
dogImages.zip  solution.py  TrainedModels
(pytorch) [root@ip-172-31-93-19 ~]# ls TrainedModels/
model.pth
(pytorch) [root@ip-172-31-93-19 ~]#
```

Step 3: Lambda function setup

After training and deploying the model, setting up a Lambda function, using the `lamdafunction.py` script, is an important task as this enables the model to be accessed by API's and other programs.

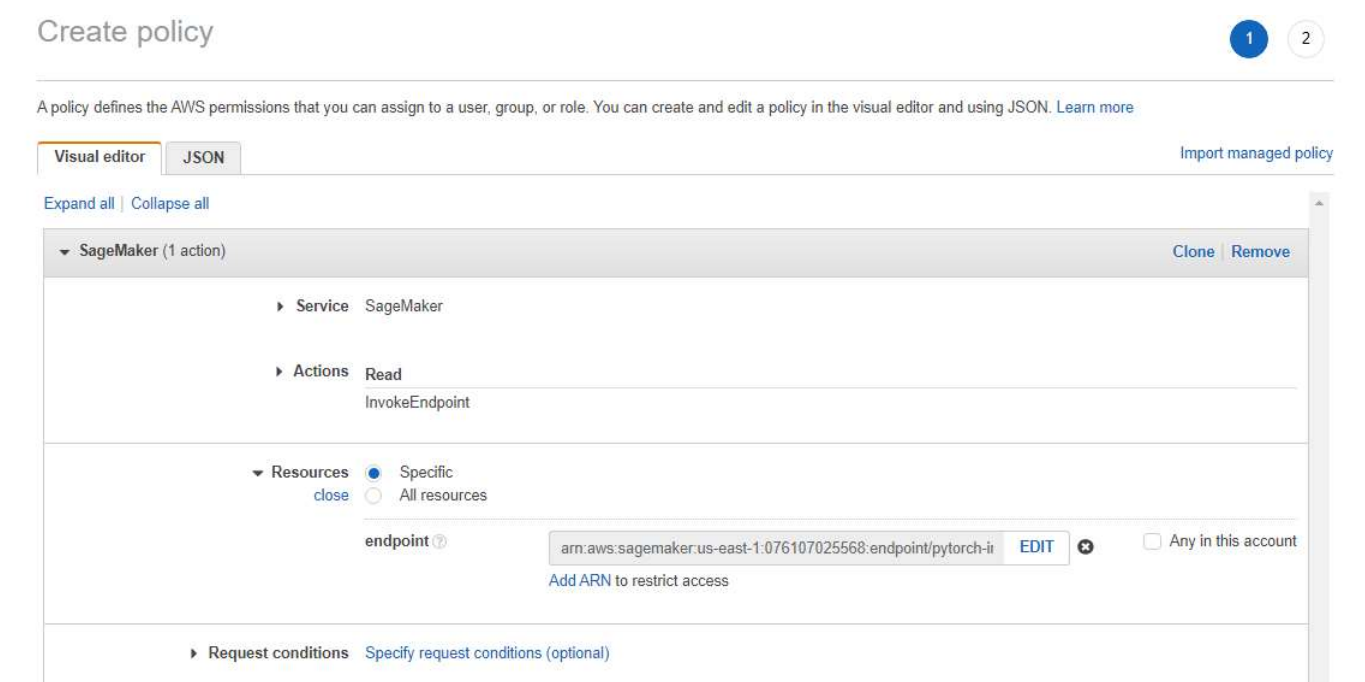
Step 4: Lambda security setup and testing

Once the Lambda function is configured, it will return an error when performing a request test to invoke the Sagemaker deployed endpoint. This error is returned since the Lambda function role has no permissions to access the endpoint.

When choosing a proper policy to resolve the security issue, it is necessary to evaluate some vulnerabilities such as the following:

- Setting up a “FullAccess” policy without restrictions is a vulnerability that can be exploited by external elements, such as cybercriminals, in case of an unauthorized access to the Lambda Function.
- Roles that are old or inactive pose a security risk that can compromise a Lambda Function, hence they should be identified and removed.
- Roles with policies for Lambda functions that are no longer in use are also a potential risk of unauthorized access and they should be revoked and deleted.

For this project, a suitable policy that limits access to invoke a specific Sagemaker endpoint is created and attached to the Lambda role using the IAM interface, as follows:



[Permissions](#) | [Trust relationships](#) | [Tags](#) | [Access Advisor](#) | [Revoke sessions](#)

Permissions policies (2) [Info](#)

You can attach up to 10 managed policies.

< 1 >

<input type="checkbox"/>	Policy name ↗	Type	Description
<input type="checkbox"/>	+ AWSLambdaBasicExecutionRole-008e7ebe-d4da-4fc0-991a-551f73c05b31	Customer managed	
<input type="checkbox"/>	+ lambdaInvokeEndpoint	Customer inline	

The results of a successful invoking test performed on the Lambda function are shown below:

Lambda > Functions > project4Inference

project4Inference

Throttle | Copy ARN | Actions

Function overview [Info](#)

Code | Test | Monitor | Configuration | Aliases | Versions

Execution result: succeeded [logs](#)

Details

The area below shows the last 4 KB of the execution log.

```
{
  "statusCode": 200,
  "headers": {
    "Content-Type": "text/plain",
    "Access-Control-Allow-Origin": ""
  },
  "type-result": "<class 'str'",
  "Content-Type-In": "<__main__.LambdaContext object at 0x7f7a3ec646a0>",
  "body": "[[0.3748210370540619, -0.046208158135414124, -0.19332340359687805, 0.18459755182266235, 0.6422909498214722, 0.35926416516304016, 0.07143840193748474, 0.2013409584760666, -0.40662631392478943, -0.2238309532403946, 0.1126127690076828, 0.2525499165058136, 0.09177562594413757, 0.09780372679233551,
```

```

{
  "statusCode": 200,
  "headers": {
    "Content-Type": "text/plain",
    "Access-Control-Allow-Origin": "*"
  },
  "type-result": "<class 'str'>",
  "Content-Type-In": "<__main__.LambdaContext object at 0x7f7a3ec646a0>",
  "body": "[[0.3748210370540619, -0.046208158135414124, -0.19332340359687805, 0.18459755182266235, 0.6422909498214722, 0.35926416516304016, 0.07143840193748474, 0.2013409584760666, -0.40662631392478943, -0.2238309532403946, 0.1126127690076828, 0.2525499165058136, 0.09177562594413757, 0.09780372679233551, 0.5082946419715881, 0.2655788064002991, 0.3439544439315796, -0.4079475402832031, -0.18330448865890503, 0.4885449707508087, 0.2703157365322113, 0.15254491567611694, 0.4578859508037567, 0.4886816740036011, -0.2267102301120758, 0.06121187284588814, 0.08620661497116089, -0.07932054996490479, 0.49667492508888245, 0.3141535520553589, 0.07674485445022583, 0.1663159877061844, -0.25284793972969055, 0.21378761529922485, 0.25132060050964355, 0.5610526204109192, 0.18995662033557892, 0.13628947734832764, 0.1684465855360031, -0.15521439909934998, 0.4581757187843323, 0.1620015799999237, 0.30257973074913025, 0.09637795388698578, 0.04541651904582977, 0.1521347463130951, -0.08743750303983688, 0.04993806779384613, -0.20915032923221588, -0.21284787356853485, -0.02543587237596512, 0.1883545070886612, -0.30492720007896423, 0.18603497743606567, -0.2546696066856384, 0.1975076049566269, 0.4556020200252533, 0.16756783425807953, 0.12390435487031937, -0.1146358922123909, 0.11059819906949997, -0.11682835221290588, 0.0710432380437851, -0.4874916970729828, 0.05493775010108948, -0.4273858070373535, -0.10279837250709534, 0.12126364558935165, 0.29550567269325256, 0.18309897184371948, 0.06595314294099808, -0.025738196447491646, -0.30459538102149963, 0.01492330338805914, -0.15622296929359436, 0.18880636990070343, -0.2251722812652588, -0.2873472273349762, 0.1890975534915924, -0.003290688619017601, -0.10139472037553787, 0.29221588373184204, 0.055993661284446716, 0.1098623126745224, -0.1766642928123474, 0.08243223279714584, 0.3377177119255066, 0.06499467045068741, 0.24551309645175934, 0.3099081516265869, -0.04492253065109253, -0.2474137842655182, -0.5685271620750427, -0.3114491403102875, -0.39963656663894653, -0.33272674679756165, -0.08196931332349777, -0.1892833411693573, -0.09226399660110474, -0.4018873870372772, -0.06535250693559647, -0.4086303412914276, -0.07361125200986862, -0.02084951475262642, -0.508001446723938, -0.1040288507938385, -0.4590185582637787, -0.9188807606697083, -0.010974306613206863, -0.4687262773513794, -0.05386451259255409, 0.1874505877494812, -0.014306633733212948, -0.3130384385585785, 0.3803076446056366, -0.5665944814682007, -0.14233903586864471, -0.19892369210720062, -0.6708981990814209, -0.24014577269554138, -0.7050748467445374, -0.48393329977989197, -0.051612406969070435, -0.26194262504577637, -0.480593204498291, -0.6064732670783997, 0.01859777607023716, -0.3616325557231903, 0.152866393327713, 0.26529568433761597, -0.5362987518310547, -0.5450462102890015, -0.34911587834358215]]]"
}

```

Step 5: Lambda concurrency setup and endpoint auto-scaling

In order to prepare a Lambda function to respond to high traffic demand with low latency, i.e multiple invocations at once, provisioned concurrency can be set up. The price depends on the amount of memory is allocated to the function and the amount of concurrency. For the purposes of this project, a small amount of

concurrency is set up, as follows:

project4Inference

ThrottleCopy ARNActions

Function overviewInfo

CodeTestMonitorConfigurationAliasesVersions

General configuration

Triggers

Permissions

Destinations

Function URL

Environment variables

Tags

VPC

Monitoring and operations tools

Concurrency

ConcurrencyEdit

Function concurrencyUnreserved account concurrency999

Provisioned concurrency configurations (1)
To enable your function to scale without fluctuations in latency, use provisioned concurrency. You can use Application Auto Scaling to automatically adjust provisioned concurrency to maintain a configured target utilization. Provisioned concurrency runs continually and has separate pricing for concurrency and execution duration. [Learn more](#)

RefreshEditRemoveAdd

Find configuration

	Qualifier	Type	Provisioned concurrency	Status	Details
	1	version	1	Ready	-

Finally, like Lambda functions use concurrency to respond to multiple requests at the same time, Sagemaker endpoints require automatic scaling to perform in a similar scenario. For this project, the auto-scaling setup with a maximun of 2 instances is shown in the following image:

Endpoint runtime settings

Update weightsUpdate instance countConfigure auto scaling

	Variant name	Current weight	Desired weight	Elastic Inference	Instance type	Current instance count	Desired instance count	Instance min - max	Automatic scaling
	AllTraffic	1	1	-	ml.m5.large	1	1	1 - 2	Yes